




## Workshop 9

### Advanced Stream Processing Pipelines

# Workshop Objectives



Learn how to work with:


1. **Batch processing:** Create automated batch processing pipelines to keep dimensional models current as source data changes (OLTP to OLAP updates)
2. **Data staging:** Implement reliable data movement using staging tables to transfer data between source and target systems safely
3. **Monitoring:** Implement audit logging tables to track successful loads and provide operational monitoring
4. **Dimension updates:** Apply Type 1 Slowly Changing Dimension (SCD) updates to overwrite old customer data with new values, maintaining only current information

# Why Batch Processing?

- Common Use Cases
  - Daily sales reports and dashboards
  - Weekly performance summaries
  - Monthly analytics refreshes

*What regular reporting or analytics tasks in your organisation currently use or might benefit from batch processing?*

*How frequently do these need to update?*




# Why Batch Processing?

- Key Benefits
  - Efficient processing of large data volumes
  - Predictable resource usage
  - Simple scheduling and monitoring

*Looking at these benefits, where in your organisation have you seen these work well or struggle?*

*For example, has scheduling regular loads ever been challenging?*



# Workshop Alignment with IFATE Pass Descriptors

Star Schemas and Data Warehousing (K15)	Building and maintaining a star schema with Type 1 SCD updates in Azure Synapse
Data Engineering Tools (K20)	Practical experience with Azure SQL, Azure Synapse, and Dedicated SQL Pools
Pipeline Deployment (K8)	End-to-end deployment of automated batch update pipeline
Data Store Monitoring (K1, S7)	Implementing audit logging and pipeline monitoring for reliable updates



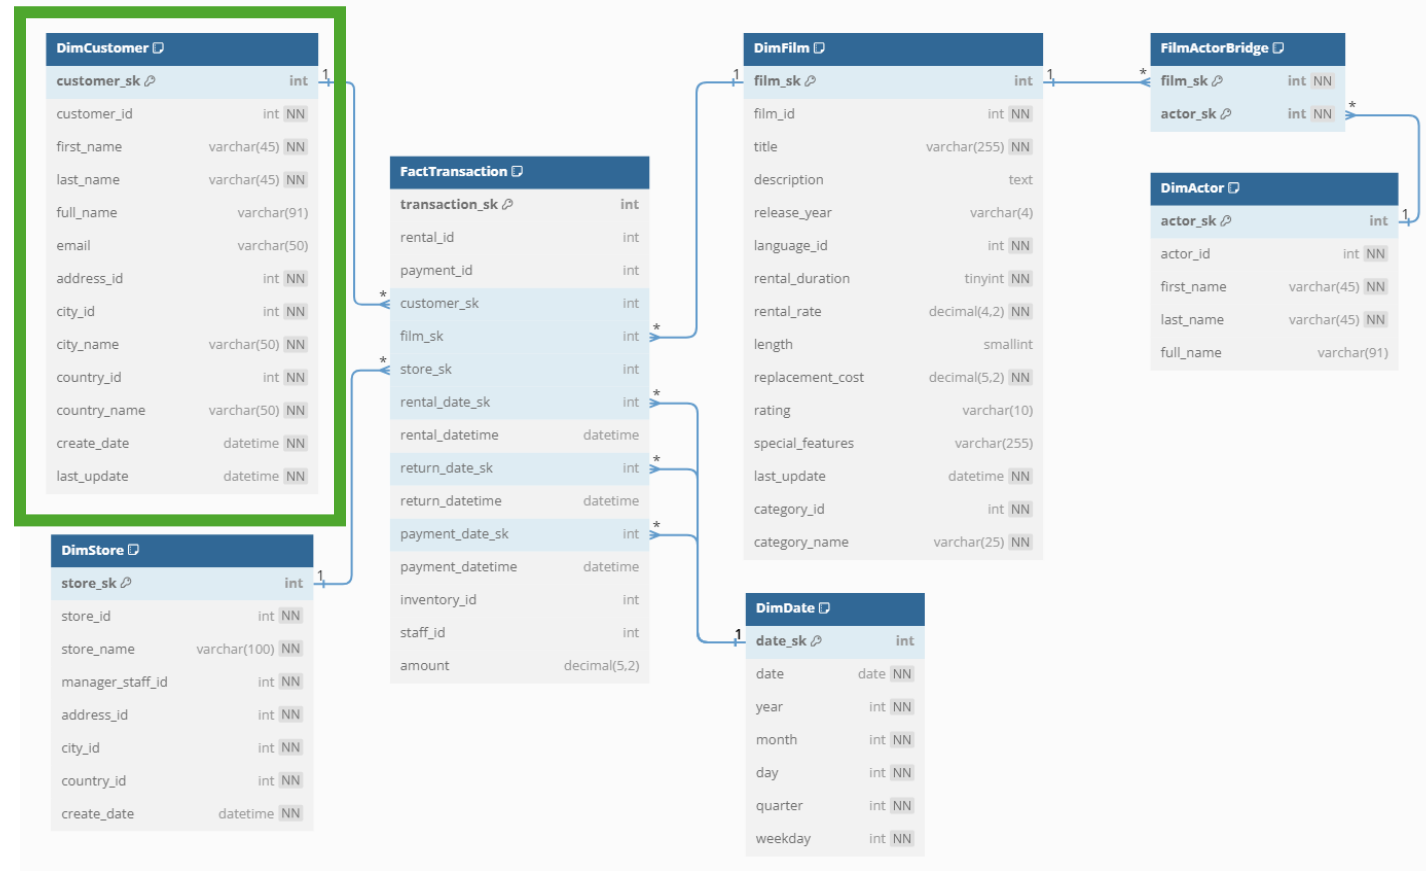
# Morning activity

Building the pipeline

# What We'll Build Today

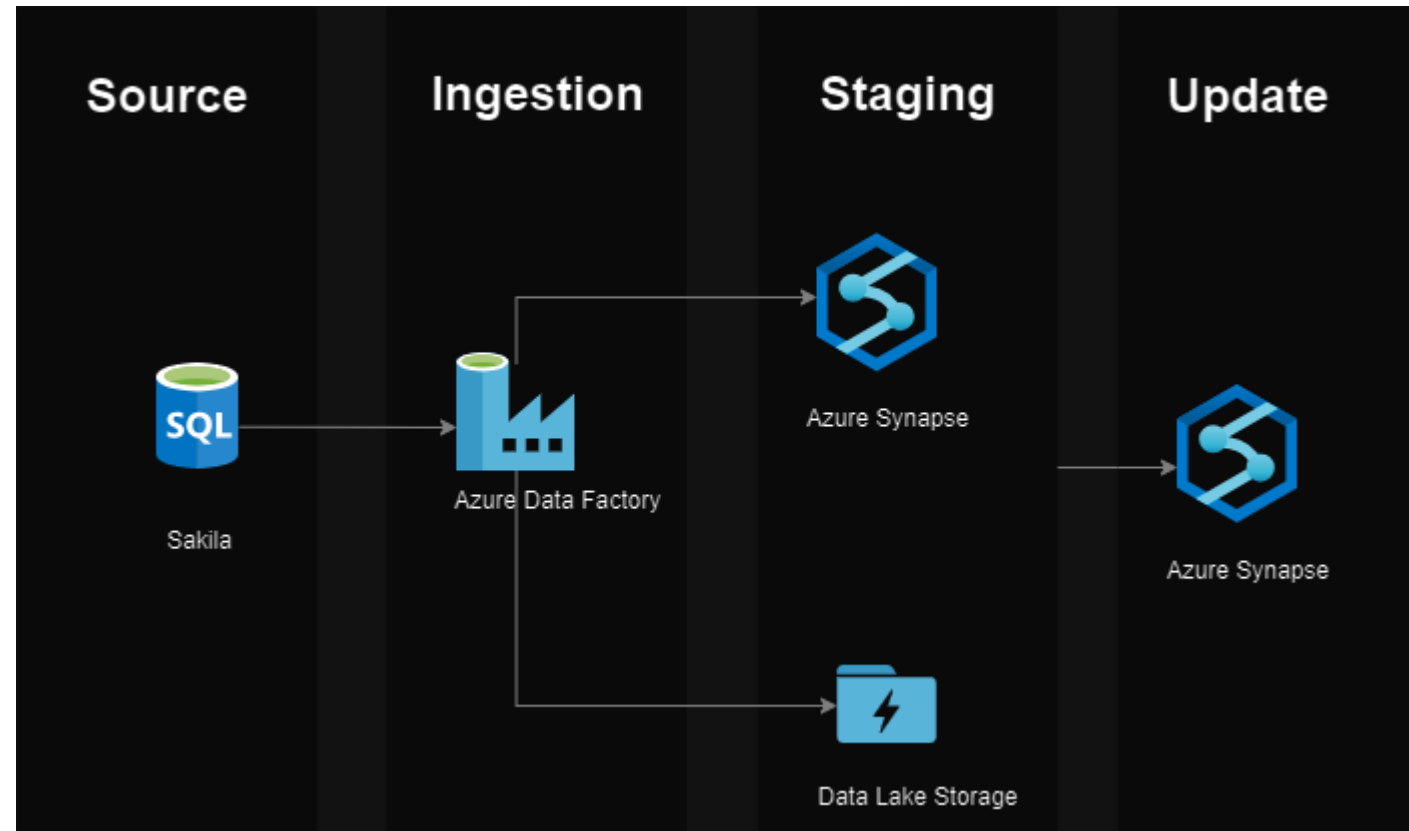
- Automated data warehouse updates using Azure Synapse
- Focus on Customer dimension maintenance
- Type 1 SCD implementation with audit logging

*Building on Workshop 3's static star schema and Workshop 6's design work*



# Our Architecture

- Key Components
  - **Source:** Sakila OLTP database in Azure SQL (OLTP)
  - **Ingestion:** Azure Data Factory pipeline
  - **Staging:** Choice of:
    - Azure Synapse staging tables
    - Data Lake Storage
  - **Update:** Azure Synapse dimension table (OLAP)
- Pipeline Operations.
  1. Extract the latest customer data
  2. Stage with audit columns
  3. Process Type 1 SCD updates





# Workshop Flow

## Morning

- Building the Pipeline
  1. Environment setup
  2. Schema and table creation
  3. Pipeline development
  4. Testing with simulated changes
- Going Further (time permitting)
  - Data Lake integration
  - Data Flow transformations
  - PySpark notebook processing

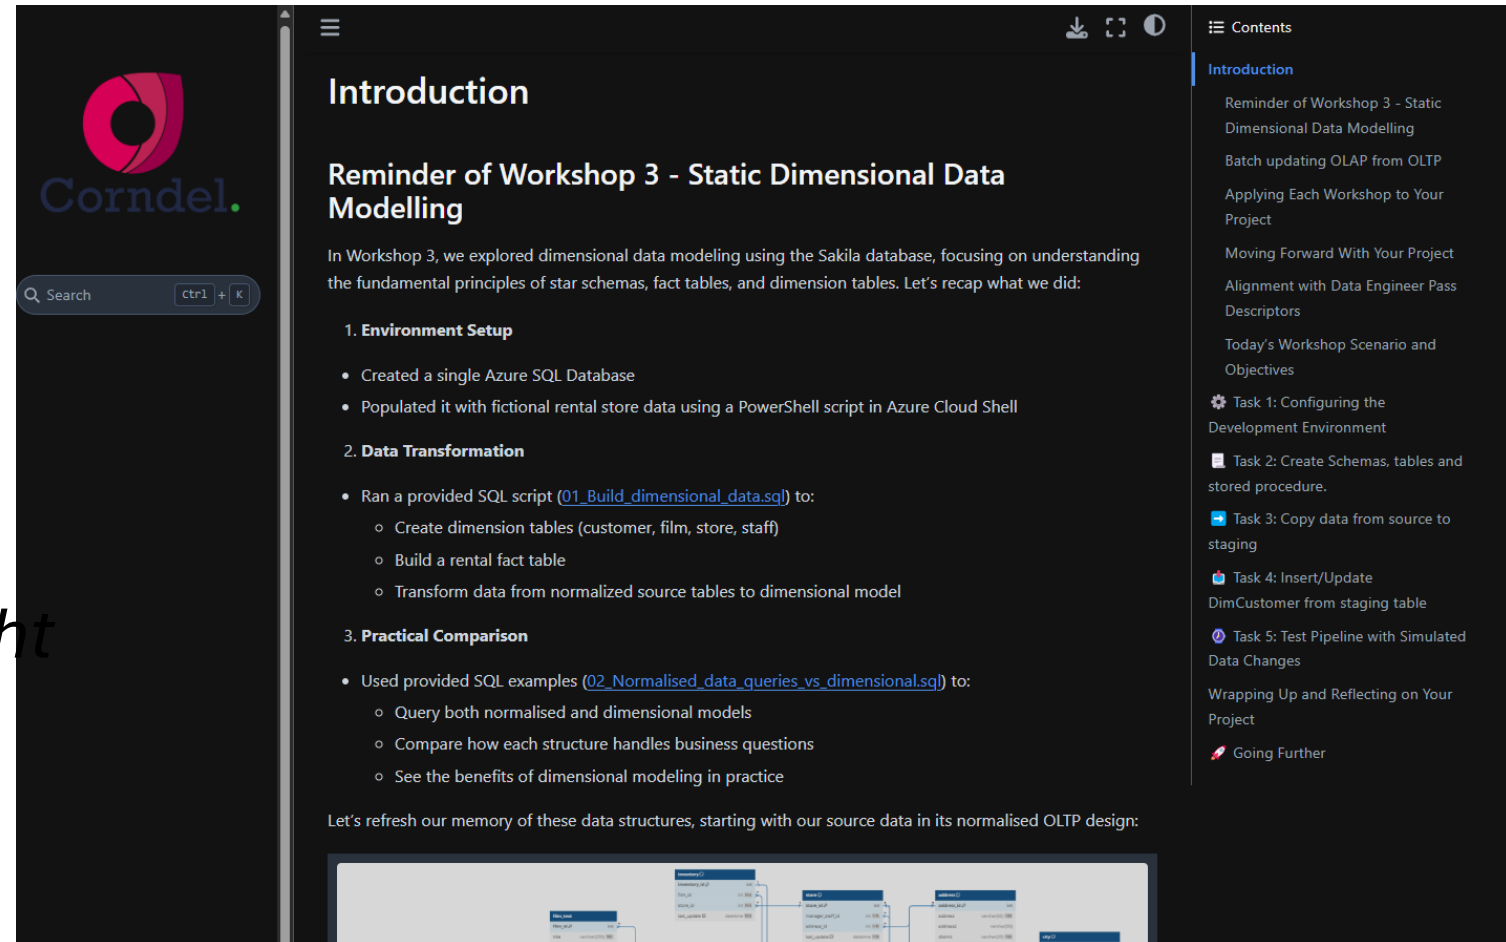
## Afternoon

- Workshop Reflection
  - What went well?
  - Technical challenges?
  - Learning points and insights?
- Project Application
  - Project data sources
  - Update patterns and data models
  - Implementation strategies

# Workshop Instructions: Your Step-by-Step Guide

- Key concepts and terminology
- Task breakdown and dependencies
- Optional "Going Further" sections
- How to ask for help

*We will now walk through the document structure and highlight how each section builds on the previous one to create a complete pipeline*





# Afternoon activity

Reflecting on the morning activity and how it applies to your project

# Workshop reflection



## Technical Implementation

- 1 ✨ Which parts of the pipeline came together smoothly?
- 2 🚧 Where did you encounter challenges?
- 3 🔍 What debugging approaches worked best?



## Learning Journey

- 1 💡 What key concepts clicked today?
- 2 🤔 Which concepts need more exploration?
- 3 ⚖️ How does this compare to any previous pipeline work?

# Build on Your Learning from E-Learning Module 06.3



- The next questions in the following two slides distil key ideas from **06.3 Batch Processing: OLTP Data Source Integration**, which you may have completed prior to this workshop.
- If you haven't yet completed it, we encourage you to do so. It provides:
  - A foundation in batch processing concepts and their practical applications.
  - Key reflective prompts to align these ideas with your project and role.
- Use questions in 6.3 to deepen your understanding of:
  - The batch processing patterns most relevant to your project.
  - Robust pipeline design, update patterns, and optimisation techniques.
  - How these align with user and business needs
- These reflections aren't just for today:
  - Use them as a starting point for future discussions with your PDE in 121s.
  - Identify areas for improvement and questions to explore further with your PDE.

# Project Application Discussion



## Project data sources

- 1 What type(s) of data sources does your project handle? 🗄️ OLTP databases? 📄 Files? 🌐 APIs?
- 2 What challenges do the formats present? ⚠️
- 3 What access methods are needed for each source type, and how are they configured securely? 🔒
- 4 Does the origin of your data (e.g., normalised tables, JSON files) impact your pipeline design? 📄
- 3 What security and access patterns do you need to consider? 🔒



## Update patterns

- 1 What update pattern best fits your use case? 🔄 Full refresh 🔄 incremental updates 📄 CDC?
- 2 🤔 What drives this choice (data volume, change frequency, etc)?
- 3 How can you track and validate changes? ✅
- 4 What could be your strategy for handling failed updates or reprocessing? 🚨

# Project Application Discussion



## Data Modelling & Requirements

- 1 Have you defined clear user and business requirements yet? 📝
- 2 Does your chosen data model support these effectively? 🎯
- 3 How do you handle historical data needs? Is Type 1 SCD sufficient or do you need Type 2? 🔄
- 4 How do you validate data quality? ✅



## Implementation and Optimisation

- 1 What tools and frameworks are you using/considering? ☁️ Cloud-native services? 💻 Local or on-prem tools? 🎵 Orchestration tools?
- 2 🤔 Why are these the right choices for your needs?
- 3 How are you optimising performance? 📊 Partitioning strategy? 🏷️ Indexing approach? 📁 Compression techniques? 🔍 Query optimisation? 💻 Hardware/resource allocation? ⚙️ Parallel processing?
- 4 What monitoring and logging is appropriate? 📈
- 5 Are your pipelines idempotent, and how do you handle retries?